

```

# Load required libraries for data manipulation, visualization, and reporting
library(knitr)          # For generating R Markdown reports
library(tidyverse)      # For data manipulation and visualization (includes dplyr, ggplot2, etc.)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.0.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(tinytex)        # For rendering R Markdown to PDF if needed

library(dplyr)          # For data manipulation (part of tidyverse, explicitly loaded for clarity)
# Install and load required packages
if (!require("dplyr")) install.packages("dplyr")
if (!require("reshape2")) install.packages("reshape2")

## Loading required package: reshape2
##
## Attaching package: 'reshape2'
##
## The following object is masked from 'package:tidyr':
##
##      smiths

library(dplyr)
library(reshape2)

# Check current working directory to ensure the file path is correct
getwd()

## [1] "C:/Users/amand/OneDrive/Desktop/Assignment2/RprojectAssignment2"

# Import the combined dataset from a CSV file
datacombined2 <- read.csv("library_data.csv", header = TRUE, sep = ",")

# Why Import This Way?
# - The dataset 'library_data.csv' contains library performance metrics across multiple years.
# - Specifying header = TRUE and sep = "," ensures the CSV is read correctly with column names.
# - This dataset includes columns like TotalOperatingRevenues, XofActiveLibraryCardholders, etc., which

# Clean column names to remove spaces and special characters for easier manipulation
colnames(datacombined2) <- gsub("[^[:alnum:]]", "", colnames(datacombined2))

```

```

# Why Clean Column Names?
# - Column names with spaces or special characters (e.g., "X of Active Library Cardholders") can cause
# - Using gsub("[^[:alnum:]]", "", ...) ensures names are alphanumeric (e.g., "XofActiveLibraryCardhold
# - This step prevents errors during data manipulation and improves code readability.

# Rename duplicate column names to avoid conflicts
dup_cols <- duplicated(colnames(datacombined2))
if (any(dup_cols)) {
  colnames(datacombined2)[dup_cols] <- paste0(colnames(datacombined2)[dup_cols], "_dup")
}

# Why Handle Duplicates?
# - Duplicate column names can cause unexpected behavior in R (e.g., during subsetting or summarization.
# - Appending "_dup" to duplicates ensures all columns are uniquely identifiable.

# Ensure key columns are in the correct format for analysis
datacombined2$Year <- as.character(datacombined2$Year)
datacombined2$Library <- as.character(datacombined2$Library)
datacombined2$TotalOperatingRevenues <- as.numeric(as.character(datacombined2$TotalOperatingRevenues))
datacombined2$XofActiveLibraryCardholders <- as.numeric(as.character(datacombined2$XofActiveLibraryCard
datacombined2$PopulationResident <- as.numeric(as.character(datacombined2$PopulationResident))
datacombined2$LocalOperatingGrant <- as.numeric(as.character(datacombined2$LocalOperatingGrant))
datacombined2$Donations <- as.numeric(as.character(datacombined2$Donations))
datacombined2$SelfgeneratedRevenue <- as.numeric(as.character(datacombined2$SelfgeneratedRevenue))
datacombined2$Staffingexpenditure <- as.numeric(as.character(datacombined2$Staffingexpenditure))
datacombined2$TotalAnnualDirectCirculation <- as.numeric(as.character(datacombined2$TotalAnnualDirectCi

## Warning: NAs introduced by coercion

datacombined2$Xofprogramsheldannually <- as.numeric(as.character(datacombined2$Xofprogramsheldannually))
datacombined2$Annualprogramattendance <- as.numeric(as.character(datacombined2$Annualprogramattendance))
datacombined2$XofPublicaccessworkstations <- as.numeric(as.character(datacombined2$XofPublicaccessworks
datacombined2$MainLibrarytotalhoursopenperweek <- as.numeric(as.character(datacombined2$MainLibrarytotal
datacombined2$ProjectGrants <- as.numeric(as.character(datacombined2$ProjectGrants))

# Why Convert Data Types?
# - Year and Library are treated as categorical variables (character type) for grouping and labeling.
# - Numeric columns (e.g., TotalOperatingRevenues) must be numeric for calculations like division or co
# - Using as.numeric(as.character(...)) handles cases where numbers might be stored as factors or text,

# Create a new column: Operating Revenue per Active Cardholder
datacombined2 <- datacombined2 %>%
  mutate(RevPerCardholder = TotalOperatingRevenues / XofActiveLibraryCardholders)

# Why Create This Column?
# - RevPerCardholder measures how much revenue each active library cardholder generates on average.
# - This metric helps assess the financial efficiency of libraries in serving their active users.
# - It's a key performance indicator (KPI) for understanding how well resources are utilized per user.

# Remove rows where RevPerCardholder is NA, infinite, or exceeds 2,500, and ensure no NA in key variabl
datacombined2 <- datacombined2 %>%
  filter(!is.na(RevPerCardholder) & is.finite(RevPerCardholder) & RevPerCardholder <= 2500 &

```

```

!is.na(TotalAnnualDirectCirculation))

# Why Filter These Rows?
# - !is.na(RevPerCardholder): Removes rows where RevPerCardholder is missing (e.g., due to missing TotalAnnualDirectCirculation)
# - is.finite(RevPerCardholder): Removes infinite values (e.g., if XofActiveLibraryCardholders is 0, can't calculate revenue per cardholder)
# - RevPerCardholder <= 2500: Removes outliers (values above 2,500 are unrealistic for revenue per cardholder)
# - !is.na(TotalAnnualDirectCirculation): Ensures no missing values in TotalAnnualDirectCirculation, a key funding source

# Why Set the Threshold at 2,500?
# - A threshold of 2,500 was chosen as a reasonable upper limit based on domain knowledge: it's highly unlikely for a library to have a revenue per cardholder exceeding 2,500
# - This threshold helps exclude data entry errors or anomalies (e.g., incorrect revenue or cardholder count)

# Calculate the correlation between Revenue per Cardholder and Local Operating Grant
insight1 <- datacombined2 %>%
  summarise(Correlation = cor(RevPerCardholder, LocalOperatingGrant, use = "complete.obs"))

# Display the correlation
print("Insight 1: Correlation between Revenue per Cardholder and Local Operating Grant")

## [1] "Insight 1: Correlation between Revenue per Cardholder and Local Operating Grant"

print(insight1)


##      Correlation
## 1      0.0118961

# Why Calculate This Correlation?
# - We want to understand how strongly LocalOperatingGrant (a key funding source) influences RevPerCardholder
# - A positive correlation suggests that more local funding leads to higher revenue efficiency per cardholder
# - The 'complete.obs' argument ensures only rows with non-missing values for both variables are used, avoiding NA values

# Bubble plot visualization to explore the relationship
ggplot(datacombined2, aes(x = LocalOperatingGrant, y = RevPerCardholder)) +
  geom_point(aes(size = PopulationResident), alpha = 0.4, color = "darkgreen") +
  geom_smooth(method = "lm", se = FALSE, color = "black") +
  labs(title = "Revenue per Cardholder vs. Local Operating Grant",
       x = "Local Operating Grant ($)",
       y = "Revenue per Cardholder ($)",
       size = "Population") +
  theme_minimal()

## 'geom_smooth()' using formula = 'y ~ x'

```



```
# Why Use a Bubble Plot?
```

```
# - The x-axis (LocalOperatingGrant) and y-axis (RevPerCardholder) show the primary relationship.
# - Bubble size (PopulationResident) adds a third dimension, showing how population size might influence revenue.
# - The linear trend line (geom_smooth) helps visualize the overall direction of the relationship.
# - Alpha = 0.4 ensures overlapping points are visible, and theme_minimal() keeps the plot clean for presentation.
```

```
# Insight 1 Interpretation:
```

```
# - A positive correlation (if observed) indicates that libraries with greater local funding achieve higher revenue per cardholder.
# - Larger bubbles (higher population) may cluster differently, suggesting that population size influences revenue.
# - For example, larger populations might have economies of scale, allowing more efficient use of funds.
```

```
# Summarize average Revenue per Cardholder by Year and Population Group
```

```
insight2 <- datacombined2 %>%
  mutate(PopulationGroup = cut(PopulationResident, breaks = quantile(PopulationResident, probs = 0:3/3,
    labels = c("Small", "Medium", "Large"), include.lowest = TRUE)) %>%
  group_by(Year, PopulationGroup) %>%
  summarise(AvgRevPerCardholder = mean(RevPerCardholder, na.rm = TRUE), .groups = "drop")
```

```
# Display the summarized data
```

```
print("Insight 2: Revenue per Cardholder Distribution by Year and Population")
```

```
## [1] "Insight 2: Revenue per Cardholder Distribution by Year and Population"
```

```
print(insight2)
```

```
## # A tibble: 15 x 3
##   Year PopulationGroup AvgRevPerCardholder
##   <chr> <fct>                <dbl>
## 1 2006 Small                186.
## 2 2006 Medium              94.3
## 3 2006 Large              96.6
## 4 2007 Small                183.
## 5 2007 Medium              99.1
## 6 2007 Large              98.4
## 7 2008 Small                171.
## 8 2008 Medium              98.6
## 9 2008 Large              105.
## 10 2009 Small              201.
## 11 2009 Medium             100.
## 12 2009 Large             110.
## 13 2010 Small              203.
## 14 2010 Medium             106.
## 15 2010 Large             114.
```

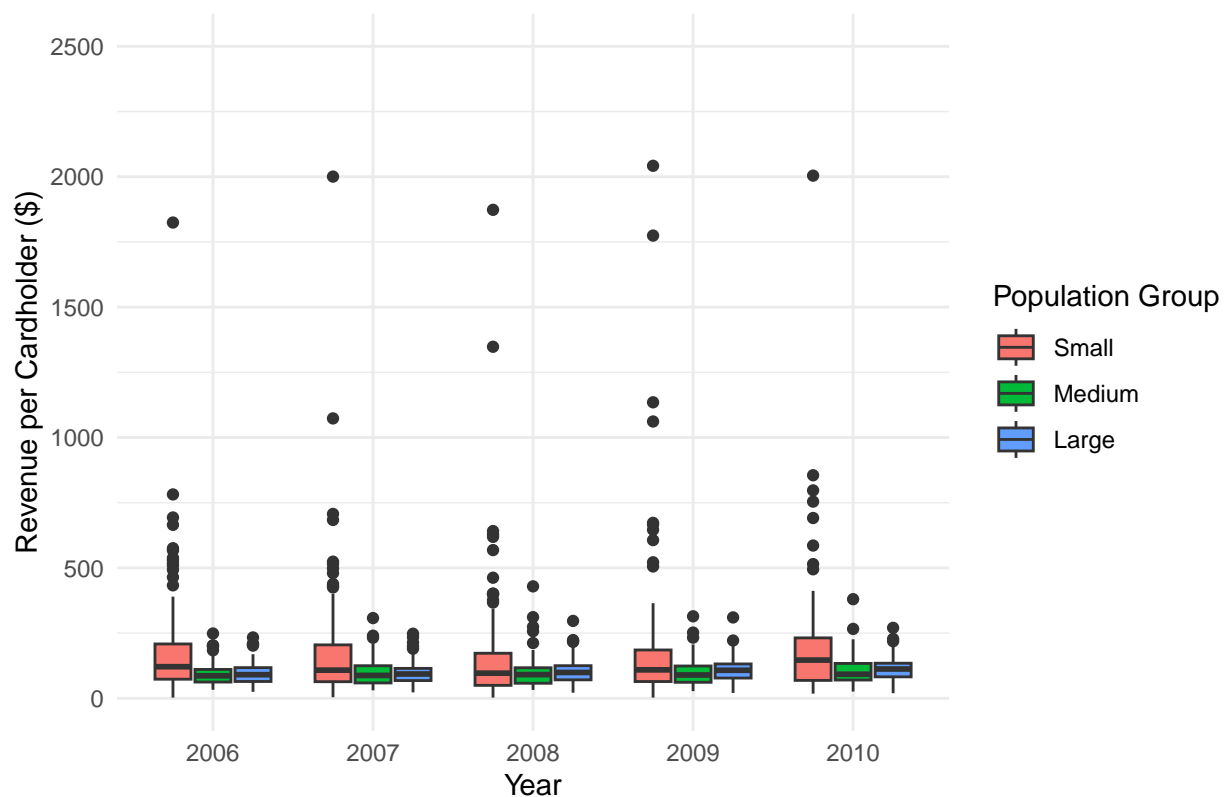
```
# Why Summarize This Way?
```

```
# - We categorize PopulationResident into Small, Medium, and Large groups using quantiles (tertiles: 0.33, 0.66, 1.0).  
# - This allows us to compare RevPerCardholder across different population sizes and years.  
# - Grouping by Year and PopulationGroup helps identify trends over time and across community sizes.
```

```
# Box plot visualization to show distribution
```

```
ggplot(datacombined2, aes(x = Year, y = RevPerCardholder,  
                          fill = cut(PopulationResident, breaks = quantile(PopulationResident, probs = c(0.33, 0.66, 1.0)),  
                          labels = c("Small", "Medium", "Large"), include.lowest = TRUE))) +  
  geom_boxplot() +  
  labs(title = "Revenue per Cardholder Distribution by Year and Population",  
        x = "Year", y = "Revenue per Cardholder ($)", fill = "Population Group") +  
  ylim(0, 2500) +  
  theme_minimal()
```

## Revenue per Cardholder Distribution by Year and Population



```
# Why Use a Box Plot?
```

```
# - Box plots show the distribution (median, quartiles, outliers) of RevPerCardholder for each Year and
# - The fill aesthetic differentiates Small, Medium, and Large population groups, making it easy to compare
# - ylim(0, 2500) ensures the y-axis aligns with our filtering threshold, keeping the plot focused on relevant
# - This visualization highlights trends over time and differences across population sizes.
```

```
# Insight 2 Interpretation:
```

```
# - If RevPerCardholder increases over time for certain population groups, it might indicate improving
# - If Small population libraries consistently have lower RevPerCardholder, they may face greater challenges
# - Outliers in the box plot might indicate specific libraries that are exceptionally efficient or inefficient
# Save the updated dataset with the new column and filtered rows
```

```
# -----
# INSIGHT 3: HEATMAP (Top 10 Libraries only)
# Shows Avg Rev per Cardholder across Libraries & Years
# -----
```

```
# Filter top 10 libraries by overall average RevPerCardholder
```

```
top_libraries <- datacombined2 %>%
  group_by(Library) %>%
  summarise(OverallAvgRev = mean(RevPerCardholder, na.rm = TRUE)) %>%
  top_n(10, OverallAvgRev) %>%
  pull(Library)
```

```

# Filter main data
filtered_data <- datacombined2 %>%
  filter(Library %in% top_libraries)

# Group and reshape
heatmap_data <- filtered_data %>%
  group_by(Library, Year) %>%
  summarise(AvgRev = mean(RevPerCardholder, na.rm = TRUE)) %>%
  ungroup()

```

## 'summarise()' has grouped output by 'Library'. You can override using the  
## '.groups' argument.

```


heatmap_matrix <- dcast(heatmap_data, Library ~ Year, value.var = "AvgRev")

# Prepare matrix
rownames(heatmap_matrix) <- heatmap_matrix$Library
heatmap_matrix <- heatmap_matrix[, -1]
heatmap_matrix[is.na(heatmap_matrix)] <- 0

# Plot heatmap
heatmap(as.matrix(heatmap_matrix),
  Rowv = NA, Colv = NA,
  col = colorRampPalette(c("white", "lightgreen", "darkgreen"))(50),
  scale = "column",
  margins = c(8, 10),
  main = "Top 10 Libraries: Revenue per Cardholder by Year")

```

## 10 Libraries: Revenue per Cardholder by Year



*# WHAT THIS HEATMAP MEANS:*

*# - Rows: Libraries (Top 10 with highest average RevPerCardholder)*

*# - Columns: Years*

*# - Colors: Darker green = higher average revenue per cardholder*

*# - White = very low or 0 revenue*

*# -----*

```
write.csv(datacombined2, "library_data_with_metrics.csv", row.names = FALSE)
```

*# Why Save the Dataset?*

*# - Saving the updated dataset ensures that the new column (RevPerCardholder) and filtered data are preserved.*

*# - This file can be used for future analyses or shared with stakeholders for transparency.*

*# - row.names = FALSE prevents adding an unnecessary index column to the CSV.*